

# Store it, maybe?

Testing Cassandra Fault Tolerance with Jepsen



Nicholas Schwartzmyer  
Insight Data Engineering Fellowship  
New York

# What's the value of this project?

- Understanding system failure is vital to becoming a good engineer.
- Distributed systems are notoriously hard to reason about, even for seasoned engineers behind industry-standard distributed databases.
- Reason better while diving deep into Cassandra internals



Also...

Data drives our companies. Losing or corrupting it is BAD!




# Jepsen



Jepsen was written by Kyle Kingbury (Aphyr)  
(<https://github.com/aphyr/jepsen>)


PICK YOUR FAVORITE ENV



SET UP A CLUSTER  
OF 5 NODES

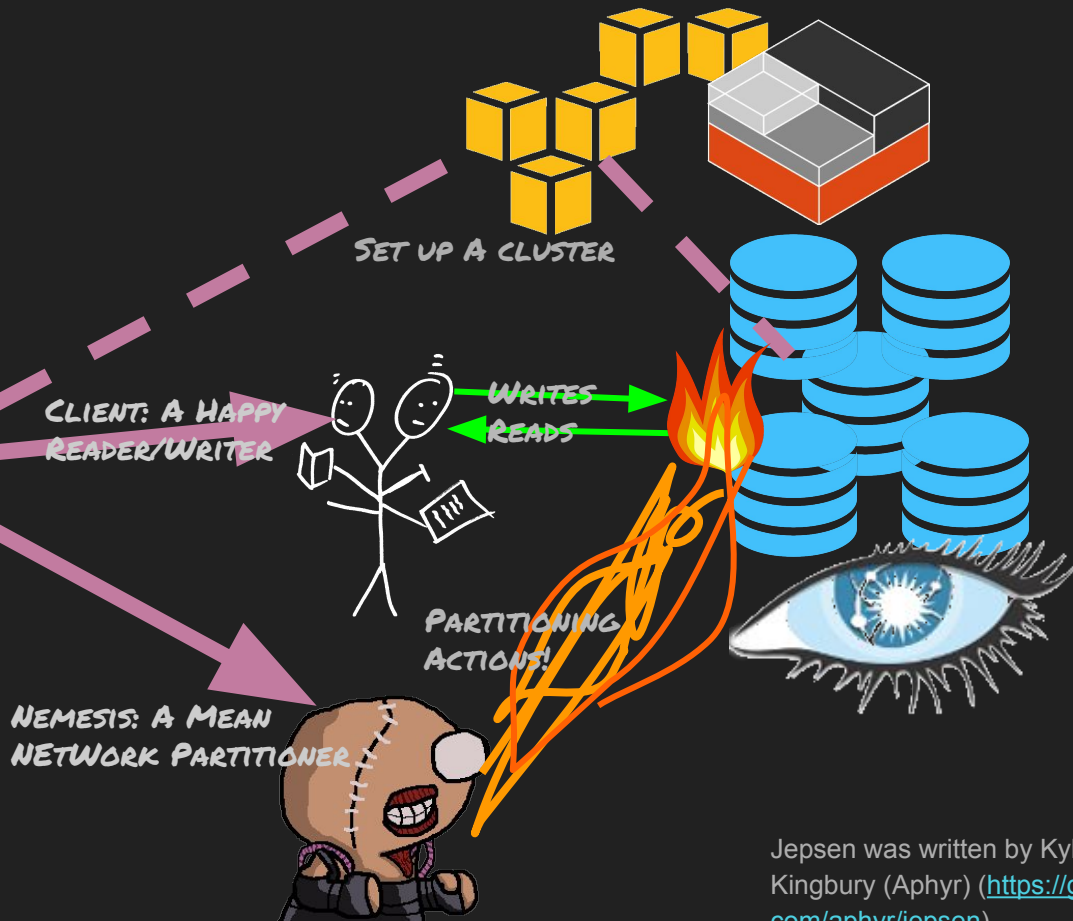


# Jepsen



Jepsen was written by Kyle Kingbury (Aphyr) (<https://github.com/aphyr/jepsen>)


# Jepsen



Jepsen was written by Kyle Kingbury (Aphyr) (<https://github.com/aphyr/jepsen>)



# Jepsen



READ BACK THE  
OP HISTORY


CHOOSE A DATASTORE!

Jepsen was written by Kyle  
Kingbury (Aphyr) (<https://github.com/aphyr/jepsen>)

# Jepsen



*MODEL: THE IDEAL SEQUENCE OF OPS*




Jepsen was written by Kyle Kingbury (Aphyr) (<https://github.com/aphyr/jepsen>)



# Jepsen



CHECK MODEL AGAINST HISTORY



Jepsen was written by Kyle Kingbury (Aphyr) (<https://github.com/aphyr/jepsen>)

# Why test Cassandra?

- Most popular wide-column datastore\*
- Emphasis on availability  $\leadsto$  *eventual consistency*
- EC  $\leadsto$  error-prone ordering decisions
- Existing tests  $\leadsto$  verifiable, incorrect  
(<https://aphyr.com/posts/294-jepsen-cassandra>)

\*<http://db-engines.com/en/ranking/wide+column+store>



# Test Plan

## Used DataStax's existing tests

(<http://www.datastax.com/dev/blog/testing-apache-cassandra-with-jepsen>)

## Cassandra 2.1.14, 2.2.6, limited 3.6

- Tests for 3.6 had underlying compatibility issues
- 3.x important when EPaxos is released

(<https://issues.apache.org/jira/browse/CASSANDRA-6246>)

## Focus on stable state tests

- Lightweight transactions
- Batch inserts
- Counters
- Set & Map operations\*

\*not tested for 3.6



# Test Plan

## Parameter tweaks:

- add/read/write consistency level  
ALL->QUORUM->TWO, SERIAL
- hinted\_handoff: (true|false)
- batch only: ATOMIC (default)/UNLOGGED  
\*UNLOGGED *not* tested by DataStax
- Keep replication factor = 3



# Results

All Results found here:

<https://github.com/nps/jepsen/tree/master/cassandra/analysis>

We'll focus on a particularly interesting case:

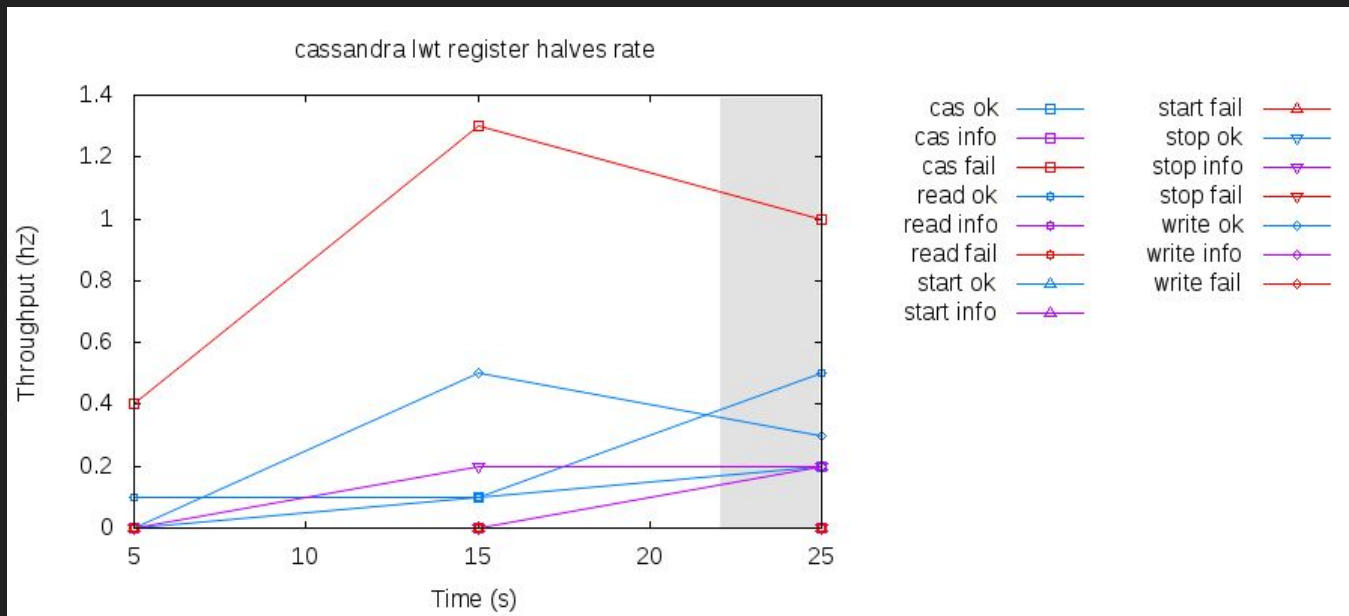
Lightweight Transactions (LWT)



# Results: Lightweight Transactions

*Ah, FAILED CAS!*

but not so fast...





# Results: Lightweight Transactions

With compare-and-set (cas), *operational* errors aren't necessarily *logical* errors!\*

Provided we can read the previous known value, we can consider it legal.

```
61 (invoke! [this test op]
62   (case (:f op)
63     :cas (try (let [[v v'] (:value op)
64                      result (cql/update conn "lwt" {:value v'}
65                                             (only-if [[= :value v]])
66                                             (where [[= :id 0]])))
67           (if (-> result first ak)
68             (assoc op :type :ok)
69             (assoc op :type :fail :value (-> result first :value))))
70   (catch UnavailableException e
```

If cas wasn't applied, mark it as failed and set it to the value of 'value'

\* See [this JIRA ticket](#) for one good discussion of the matter

# Results: Lightweight Transactions

Ensure the failed cas has a valid previous value and map it as a read success

```
40 ; A failure; fill in either value.
41 :fail
42 (let [i      (get index (:process op))
43       _      (assert i)
44       invocation (nth history i)
45       value      (or (:value invocation) (:value op))
46       [invocation' op'] (if (and (= :cas (:f op)) (number? (:value op)))
47                             [(assoc invocation :f :read :value (:value op))
48                              (assoc op :f :read :type :ok)]
49                             [(assoc invocation :value value) (assoc op :value value)]])
```


\* See [this JIRA ticket](#) for one good discussion of the matter

# Challenges!

- Understanding Jepsen mechanics by reading through the code
- Picking up some Clojure to do so!
- Significant merge conflict and versioning issues with the tests
- Crash course in Cassandra fault tolerance model
- Researching distributed data structures and state models



*KYLE KINGSBURY, JEPSEN AUTHOR*



*LESLIE LAMPORT, DISTRIBUTED SYSTEM BOSS*

# About Me

Nick Schwartzmyer

- MS, Computational Linguistics
- 8 years professional experience

*I LIKE HUNTING BUGS!*



*INCREASINGLY INTERESTED  
IN SOFTWARE CORRECTNESS*



*COMMAND LINE GEEK*





# Test Plan

Nemesis types:

- Bridge partitions
- Random node isolations
- Clock skew
- Kill a node





# Results: Atomic Batch

W=Q,R=ALL,+/-hinted\_handoff:

```
INFO jepsen.core - Everything looks good! \('-'`)/
```





Fig1: qa, +hh




W=Q,R=Q,+/-hinted\_handoff:

```
INFO jepsen.core - Everything looks good! \('-'`)/
```

Fig2: qq, -hh

# Results: Atomic Batch



```
INFO jepsen.core - Everything looks good! \('-'`)/
```

Fig1: qq, -hh

# Results: Unlogged Batch




Fig: W=Q,R=Q,-  
hinted\_handoff



Inconsistencies are possible if client & coordinator both suffer failures.

```
INFO jepsen.core - Everything looks good! \('-'`)/
```

# Results: Counters

Monotonic INC & INC/DEC tests:


```
INFO jepsen.core - Everything looks good! \('-'`)/
```



# Results: Map & Set tests

```
INFO jepsen.core - Everything looks good! \('-'`)/
```

```
{:set
 {:valid? true,
  :lost "#{}",
  :recovered "#{}",
  :ok "#{0..3525}",
  :recovered-frac 0,
  :unexpected-frac 0,
  :unexpected "#{}",
  :lost-frac 0,
  :ok-frac 1},
 :perf
 {:latency-graph {:valid? true},
  :rate-graph {:valid? true},
  :valid? true},
 :valid? true}
results.edn (END)
```



# In Summary

The original Jepsen tests  
succeeded in making the  
C\* folks better at designing for failure

Testing makes our systems better

But finding nothing wrong is also kind of sad.





# Future directions

Get tests working for C\* 3.6

Compare with Riak

- Also Dynamo-based, old Jepsen test exist
- Richer CRDTs; what's the performance hit when partitioned?

Write my own tests!

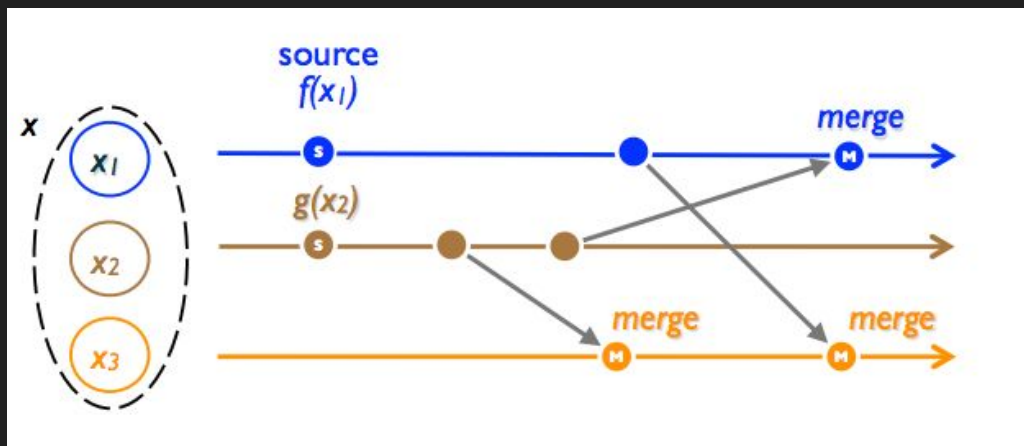
- HBase: Aphy: "I wouldn't be at all surprised if HBase is terrible haha."
- Accumulo: Growing in popularity, NSA backed

(These are backed by Zookeeper, so already have known minimum constraints, though...)



# In an even more distant future...

For fun, Implement a class of CRDT atop a datastore



Analyze how they impact with Jepsen tests

# Challenges!

The realities-of-working-with-software kind

- Getting Jepsen up-and-running
- Significant merge conflict and versioning issues with the tests

